DGX Stack OCR Test Document

Page 1 — Introduction

This document is a synthetic three-page PDF used to verify the end-to-end OCR pipeline of the DGX Stack. It exercises plain prose, a structured table, and a short technical listing so you can check whether the vision model is correctly extracting headings, paragraphs, tabular data, and code-like text.

The DGX Spark is a compact Grace-Blackwell workstation with 128 GB of unified memory, a GB10 SoC, and CUDA 13. This stack runs two models side by side: Qwen 3.6 35B NVFP4 serving the chat endpoint, and the dots.mocr document-parsing model serving the OCR endpoint.

If the OCR output below preserves the page markers, the heading structure, and the numeric values in the table on page 2, the pipeline is healthy.

Page 2 — Benchmark Table

The following table summarises hypothetical throughput numbers. The actual values are unimportant — what matters is that the OCR correctly reads every cell and preserves the column alignment.

Model	Params	Active	Weights	Context	Tok/s
Qwen 3.6 35B	35B	3B	23 GB (NVFP4)	262K	71.3
Qwen 3.5 35B	35B	3B	35 GB (FP8)	262K	58.2
Llama 3.1 70B	70B	70B	140 GB (BF16)	128K	12.4
Phi-4 14B	14B	14B	28 GB (BF16)	16K	94.6

Note: tokens-per-second numbers above are illustrative only. Real throughput depends heavily on batch size, prompt length, KV cache dtype, and whether prefix caching is enabled.

Page 3 — Example Request

Below is a representative JSON payload sent to the OpenAI-compatible chat endpoint. The OCR should preserve the braces, quotes, and the field ordering even though this is rendered as body text rather than a monospace block.


```
{
  "model": "qwen3.6-35b",
  "messages": [
    {"role": "user", "content": "Summarise page 2."}
  ],
  "max_tokens": 256,
  "temperature": 0.2
}
```


A successful OCR pass should return markdown in which page 1 contains the introduction, page 2 contains the benchmark table, and page 3 contains this example request. The word END-OF-TEST-DOCUMENT appears once, right here, so you can grep for it to confirm the final page was reached.

END-OF-TEST-DOCUMENT